RAPHAEL SAMUEL

Full-Stack Engineer | AI Agent Tooling & Data Infrastructure

samuelraphael925@gmail.com | var-raphael.vercel.app | github.com/var-raphael | Remote (UTC+1)

SUMMARY

Full-stack engineer with 7+ years across Go, TypeScript, and Python, specializing in AI-agent-native data infrastructure. Founder of three live, self-built products, spanning versioned data extraction, verified admissions data, and GitHub-based candidate verification, with native MCP servers for direct AI agent access, all designed, built, deployed, and operated solo. Comfortable owning a system end to end: backend, frontend, and the infrastructure keeping it online.

SKILLS

Languages: Go, TypeScript, Python, JavaScript
Frameworks & Libraries: Next.js, React, Tailwind CSS, LangChain
Focus Areas: AI/LLM integration, MCP servers, data pipelines & versioning, web extraction & automation, API design
Databases: PostgreSQL, MySQL, SQLite, Supabase, vector databases
Practices: REST/GraphQL APIs, schema-first data modeling, agent-facing tooling, solo end-to-end product delivery
Tools: Docker, Vercel, pandas, Matplotlib, Claude Code, Git, Paystack API, Stripe API
AI/LLM Platforms: OpenAI, Anthropic (Claude), Mistral
Web Scraping & Automation: Firecrawl, Playwright, Chromium, Crawlee, Apify

PROFESSIONAL EXPERIENCE

Founder & Solo Engineer, Quorel *2025 to Present*

quorel.vercel.app, versioned data extraction platform with native AI agent access

- Built a schema-first extraction pipeline that turns any public website into a structured, versioned dataset from a plain-English field description
- Shipped a native MCP server on every dataset, letting Claude or any MCP-compatible agent query the data, catch changes since last refresh, and publish corrections as new reversible versions
- Grew to 10 signups and 4 live public datasets covering 844+ records, serving 300+ API calls to date, with nightly and on-demand refreshes triggered by users via a cron job ping

Founder & Solo Engineer, Groundtruth *2026 to Present*

ground-truth-app.vercel.app, GitHub-based candidate verification for hiring teams

- Built a hiring tool that reads a candidate's GitHub (repos, commit history, live deployed projects) and produces a ranked, evidence-backed report scored per role, replacing resume screening
- Designed reports so every claim traces to a source: dead links are flagged rather than scored as real, and facts are kept separate from opinions
- Engineered the pipeline in Go: ingests GitHub repo and commit data, verifies live project URLs, and runs Mistral LLM analysis against each role's requirements, persisting evidence-linked results in PostgreSQL and serving reports through a Next.js/TypeScript frontend

Founder & Solo Engineer, VarsityLine *2025 to Present*

Verified Nigerian university admissions data platform

- Built a verified admissions data platform covering 24 Nigerian universities, cross-checked against each school's admissions office with timestamped "last confirmed" records
- Implemented Paystack-powered subscription billing plus search, filter, and export functionality
- Connected VarsityLine to Quorel via a custom MCP server, letting Claude Desktop autonomously monitor, clean, and push validated updates across all 24 universities with zero manual coding or scraping
- Grew to 13 paying subscribers and 132 free users, powering 687+ searches platform-wide

Freelance Full-Stack Developer, Crypto-Stock Trading Platform (Startup client, NDA) *Mar 2026*

Remote, 3-person team

- Built the full frontend and backend (Next.js, TypeScript) as part of a 3-person team, enabling users to buy stocks using Ethereum
- Owned all backend logic and handed off clean integration points to the team's Web3 engineer for on-chain settlement
- Delivered a working, client-facing platform within a 4-week engagement

OPEN SOURCE

Gnat, Self-hosted, single-binary Go analytics platform *2026*

github.com/var-raphael/gnat

- Built a self-hosted analytics platform as a single Go binary with SQLite, Postgres, or MySQL support: no Docker, no multi-service stack
- Designed a built-in MCP server so AI agents can query analytics data directly, with zero-dependency, environment-variable-only configuration